

Virtualizing FPGAs in the Cloud

ASPLOS'20

Yue Zha

Jing Li

Presenter: 黃明瀧

Outline

- ▶ Introduction
 - ▶ Objectives
 - ▶ Existing Solutions
- ▶ Methods
 - ▶ System Overview
 - ▶ Programming, Architecture Compilation, and System Layer
- ▶ Evaluation
- ▶ Conclusion

Introduction

Objectives

Simultaneously solve 2 problems:

1. Bad utilization

- ▶ Cloud providers currently allocate **physical FPGA** for user app
- ▶ Causes **internal fragmentation**

2. Scaling up is difficult (for software devs)

- ▶ Vendor tools focuses on single-device programming
- ▶ Devs have to:
 - ▶ Manually partition app into multiple FPGAs
 - ▶ Manually handle inter-FPGA communication

Existing Solutions

Problems to solve: (1) Bad utilization and (2) Difficult to use multi-FPGA

- ▶ The overlay architecture
 - ▶ Create “abstract FPGA” on top of different FPGAs
 - ▶ Analogous to JVM for software
 - ▶ Does not solve (1)
- ▶ Slot-based method
 - ▶ Analogous to MIG for GPU
 - ▶ Does not solve (2)
- ▶ AmorphOS
 - ▶ 2 modes: slot-based and combined-compile
 - ▶ Does not solve (2)

Methods

Overview of ViTAL

- ▶ Carefully divide FPGA into identical physical blocks
- ▶ Compile apps into virtual blocks
- ▶ Map virtual blocks to physical blocks at runtime
- ▶ **Runtime allocation** \Rightarrow Good utilization
- ▶ Illusion of **single, infinitely large FPGA** \Rightarrow Good scale-out

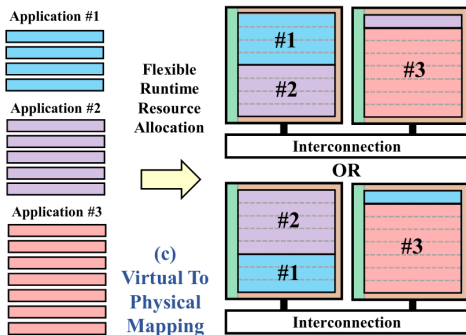


Figure 1: Virt-to-phys mapping

Programming Model

- ▶ Reuse HLS tool from Xilinx
- ▶ Programming model: **single, large FPGA**
- ▶ “Portable across FPGAs”¹. -> I don't see that true

Architecture Layer

- ▶ **Homogeneous abstraction:** “User region” blocks are the same
- ▶ “Service” and “communication” region for interconnect




Figure 2: Architecture details

Compilation Layer

1. Synthesis (HLS \rightarrow RTL \rightarrow Netlist)
2. **Partition**
 - ▶ Into virtual blocks
 - ▶ Done at the **netlist level** for **max resource usage**
3. Communication interface generation
4. P&R and Relocation (at runtime)

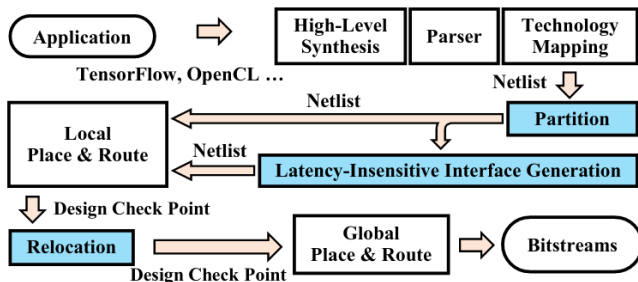


Figure 3: Compilation flow

System Layer

- ▶ Uses *partial reconfiguration* to deploy virtual blocks
- ▶ Tries to allocate **single FPGA** for one app
 - ▶ ...then 2 FPGAs, then 3 FPGAs, ...
 - ▶ Lowers communication overhead within app
- ▶ DRAM access is virtualized & managed (for isolation)




Figure 4: System controller

Evaluation

Conclusion