

Virtualizing FPGAs in the Cloud

ASPLOS'20

Yue Zha

Jing Li

Presenter: 黃明瀧

2020

Outline

- ▶ Introduction
 - ▶ Objectives
 - ▶ Existing Solutions
- ▶ Methods
 - ▶ System Overview
 - ▶ Programming, Architecture Compilation, and System Layer
- ▶ Evaluation
- ▶ Conclusion

Introduction

Objectives

Simultaneously solve 2 problems:

1. Bad utilization

- ▶ Cloud providers currently allocate **physical FPGA** for user app
- ▶ Causes **internal fragmentation**

2. Scaling up is difficult (for software devs)

- ▶ Vendor tools focuses on single-device programming
- ▶ Devs have to:
 - ▶ Manually partition app into multiple FPGAs
 - ▶ Manually handle inter-FPGA communication

Existing Solutions

Problems to solve: (1) Bad utilization (2) Difficult to use multi-FPGA

- ▶ The overlay architecture
 - ▶ Create “abstract FPGA” on top of different FPGAs
 - ▶ Analogous to JVM for software
 - ▶ Does not solve (1)
- ▶ Slot-based method
 - ▶ Analogous to MIG for GPU
 - ▶ Does not solve (2)
- ▶ AmorphOS
 - ▶ 2 modes: slot-based and combined-compile
 - ▶ Does not solve (2)

Methods

Overview of ViTAL (Virtualization Stack for FPGAs in the Cloud)

- ▶ Carefully divide FPGA into identical physical blocks
- ▶ Compile apps into virtual blocks
- ▶ Map virtual blocks to physical blocks at runtime
- ▶ **Runtime allocation** \Rightarrow Good utilization
- ▶ Illusion of **single, infinitely large FPGA** \Rightarrow Good scale-out

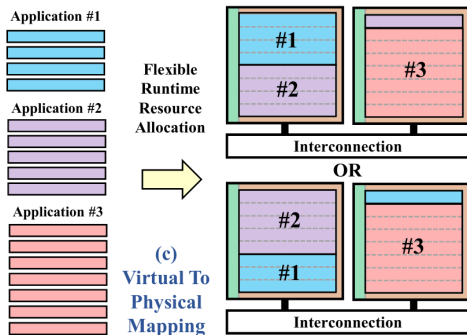


Figure 1: Virt-to-phys mapping

Programming Model

- ▶ **Single, infinitely large FPGA**
- ▶ “Portable across FPGAs”
- ▶ Reduces programming complexity
 - ▶ ViTAL takes care of partitioning & communication
 - ▶ Supports C/C++/OpenCL/other DSLs (by using HLS)

Architecture Layer

- ▶ **Homogeneous abstraction:** *User region* blocks are the same
- ▶ *Service and communication region* for interconnection




Figure 2: Architecture details

Compilation Layer

1. Synthesis (HLS \rightarrow RTL \rightarrow Netlist)
2. **Partition**
 - ▶ Into virtual blocks
 - ▶ Done at the **netlist level** for **max resource usage**
3. Communication interface generation
4. P&R and Relocation (at runtime)

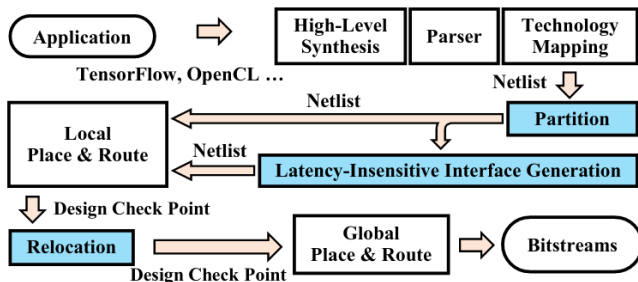


Figure 3: Compilation flow

System Layer

- ▶ Uses *partial reconfiguration* to deploy virtual blocks
- ▶ Tries to allocate blocks from **single FPGA** for one app
 - ▶ ...then 2 FPGAs, then 3 FPGAs, ...
 - ▶ Lowers communication overhead within app
- ▶ DRAM access is virtualized & managed (for isolation)




Figure 4: System controller

Evaluation

Evaluation


3 sets of benchmarks

1. Small synthetic task

- ▶ Evaluates bandwidth in *architecture layer*
- ▶ 94 Gb/s (inter-FPGA), 312.5 Gb/s (inter-die)
- ▶ Resource usage of communication region $< 10\%$

2. Deep learning task

- ▶ Evaluates *compilation & relocation time*



Evaluation (cont'd)

3 sets of benchmarks

3. Multiple deep learning tasks

- ▶ Tasks come at random time
- ▶ Simulates cloud environment
- ▶ 82% better than not sharing at all
- ▶ 25% better than AmorphOS (\because multi-FPGA)




Figure 5: Response time comparison

Conclusion

Conclusion

- ▶ Decouple **compilation** and **resource allocation** for FPGA
- ▶ Provides **single large FPGA** illusion to devs
- ▶ Achieves higher utilization & lower response time
- ▶ My thoughts
 - ▶ The claimed “portability” is false
 - ▶ Portability can be solved by **combining overlay architecture** on top of ViTAL
 - ▶ The system is applicable to non-FPGA accelerators

References

References

- ▶ Yue Zha and Jing Li. 2020. Virtualizing FPGAs in the Cloud. Proceedings of the Twenty-Fifth International Conference on Architectural Support for Programming Languages and Operating Systems. Association for Computing Machinery, New York, NY, USA, 845–858.
<https://doi.org/10.1145/3373376.3378491>
- ▶ Quraishi, Masudul Hassan and Tavakoli, Erfan Bank and Ren, Fengbo. 2020. A Survey of System Architectures and Techniques for FPGA Virtualization.